

Evaluation-Driven Representation Learning: Content Predicts Collective Reader Reactions

Zhi Liu

Independent Researcher, China

ailiheizi@gmail.com

Abstract

Do the reactions a piece of content provokes reflect learnable properties of the content itself, or are they dominated by audience culture and noise? We study this using web-novel paragraphs paired with reader paragraph-level reviews (起点中文网 danmaku-style comments). We first show a naive approach fails: classifying paragraphs into predefined reaction categories performs at chance, because ~63% of reviews are jokes/memes that swamp content signal. However, an anchor test reveals the signal *is* present: reviews of the same paragraph are significantly more similar to each other than to reviews of other paragraphs (Cohen’s $d=0.60$, $p<10^{-222}$). We then show the correct formulation—**contrastive content** ↔ **review alignment**—successfully extracts this signal: a trained two-tower model retrieves a paragraph’s true review-set on held-out chapters at $\text{Recall}@1=0.35$, 17× above random and 2.4× above the raw-embedding baseline. Our contribution is methodological: **collective evaluation is a learnable function of content, but only under a retrieval formulation that bypasses noisy categorical labels.**

1. Introduction

When readers comment on a passage of fiction, are their reactions determined by the content, or by factors external to it (memes, community culture, individual mood)? This question bears on a broader hypothesis we call **Evaluation-Driven Representation Learning (EDRL)**: that naturally-occurring human evaluations carry enough information to shape meaningful content representations, without explicit feature engineering.

We test EDRL on a clean data source—web-novel paragraphs with paragraph-level reader reviews—where content (X) and evaluation (Y) are naturally distinct. Our investigation yields a nuanced answer with a clear methodological lesson.

Contributions: 1. We show that the naive EDRL formulation (predict predefined reaction categories from content) **fails at chance level**, because meme/joke reviews dominate (63%) and categorical labels compress away the signal. 2. We establish via an anchor test that content-driven signal nonetheless exists: same-paragraph reviews cluster significantly ($d=0.60$, $p<10^{-222}$). 3. We show the **retrieval formulation succeeds**: contrastive content ↔ review alignment retrieves true review-sets at $\text{R}@1=0.35$ (17× random) on held-out chapters, beating the raw-embedding baseline. 4. We identify the key technique for small-data regimes: residual projection with early stopping, preventing the overfitting that destroys naive contrastive training.

2. Related Work

Aspect-based sentiment analysis [1] extracts evaluative aspects but requires predefined categories or labels. **Danmaku/comment analysis** [2] treats reactions as scalar sentiment, missing the content-reaction mapping. **RLHF** [3] uses curated preferences; we use natural, unsolicited reactions. **Cross-modal retrieval** (CLIP-style) [4] aligns paired modalities; we adapt this to align content with its collective evaluation. Unlike prior work, we (a) test whether reactions are a *function* of content, and (b) show a retrieval formulation succeeds where classification fails.

3. Data

Source: 起点中文网 (Qidian), novel 《圣墟》 (Shengxu), paragraph-level reader reviews scraped via the public review API. Each paragraph has its text and the reviews readers attached to it.

Scale: 30 chapters, 279 paragraphs with ≥ 2 reviews (242 with ≥ 3), 1,763 total reviews.

Key property: Content (paragraph text) and evaluation (reviews) are distinct objects—unlike video danmaku where the comment text *is* the signal, here we can ask whether content predicts reaction.

4. Experiments

4.1 Naive Formulation Fails

We label each paragraph’s dominant reaction type (via LLM aggregation over its reviews) into categories {joke, awe, emotion, analysis, confusion} and train a classifier from paragraph content.

Result: Accuracy 0.667 vs. majority baseline 0.700—**below majority**. The distribution is dominated by joke (63%), with awe/emotion nearly absent ($< 2\%$ each). Predefined categories compress the signal and meme-reviews swamp it.

Finding 1: Predicting predefined reaction categories from content fails at chance. This is the formulation most prior EDRL-style attempts would use—and it does not work.

4.2 Anchor Test: The Signal Exists

Before concluding no signal exists, we test the prerequisite directly: are reviews of the *same* paragraph more similar than reviews of *different* paragraphs?

Comparison	Mean similarity	n pairs
Same-paragraph reviews	0.417	5,947
Different-paragraph reviews	0.354	5,906

Finding 2: Same-paragraph reviews are significantly more similar ($\Delta = +0.062$, Cohen’s $d = 0.60$, $t = 32.6$, $p < 10^{-222}$). Reviews cluster by paragraph—**content does drive reaction**. The naive failure (§4.1) was a formulation problem, not an absence of signal.

4.3 Retrieval Formulation Succeeds

We train a two-tower contrastive model: $f(\text{content})$ and $g(\text{reviews})$ projected to a shared space, with a CLIP-style loss matching each paragraph’s content to its own review-set. Evaluation: on held-out chapters, rank all candidate review-sets by similarity to $f(\text{content})$; measure Recall@k. We use residual projections (output = $\text{normalize}(x + \alpha \cdot \Delta(x))$) with dropout, weight decay, and early stopping on a validation split—critical to prevent overfitting on small data.

Method	R@1	R@5	R@10
Random	0.020	0.102	0.204
Raw BGE (no training)	0.143	0.612	0.653
Learned (ours)	0.347	0.633	0.694

Finding 3: The trained model retrieves the correct review-set at $R@1 = 0.35$ — $17\times$ above random and $2.4\times$ above raw BGE. Training extracts content features that predict collective reaction. (Split by chapter prevents leakage: test chapters are entirely unseen.)

Finding 4: Naive contrastive training (plain MLP towers, no regularization) overfits catastrophically ($R@1$ drops to 0.10 below raw BGE). Residual projection + early stopping is essential in the small-data regime.

5. Discussion

Collective evaluation is a learnable function of content—but the formulation matters. Three formulations, three outcomes: - Categorical classification: **fails** (chance) — noise-dominated, compressed labels. - Raw embedding retrieval: **partial** ($R@1 = 0.14$)

— signal present but not isolated. - Contrastive alignment: **succeeds** ($R@1=0.35$) — extracts content→reaction mapping.

Why the noise doesn't kill retrieval. Even with 63% joke reviews, the aggregate review embedding retains a content-correlated component (anchor test). Classification forces a hard category decision dominated by the majority (jokes); retrieval uses the full similarity structure, where the content-correlated component still ranks the true pairing highest.

Relation to a broader boundary. This complements findings that *implicit* knowledge is hard to access [author's other work]: here, collective evaluation—seemingly noisy and subjective—turns out to carry an *explicit*, retrievable content signal, once the task avoids lossy categorical bottlenecks.

6. Limitations

1. **Single book/platform.** One novel (《圣墟》) on one platform (Qidian). Generalization to other genres, platforms, and languages is untested.
2. **Scale.** 242 paragraphs; retrieval $R@k$ is measured over tens of held-out candidates, not thousands. Larger-scale retrieval would test robustness.
3. **Aggregate reviews.** We mean-pool review embeddings; richer aggregation (attention, set encoders) may capture more.
4. **Correlation, not mechanism.** We show content predicts reaction, not *which* content features drive it. Interpretability is future work.

7. Conclusion

We show that collective reader reactions to fiction are a learnable function of content—but only under a retrieval formulation. Naive categorical prediction fails at chance because meme-reviews dominate; yet an anchor test proves the signal exists ($d=0.60$), and contrastive content↔review alignment extracts it ($R@1=0.35$, $17\times$ random). The methodological lesson generalizes: when evaluations are noisy, retrieval-based alignment recovers content-evaluation structure that categorical classification destroys.

8. Ethics Statement

All reviews are publicly posted reader comments collected via a public API, containing only text—no user identifiers or PII. Data is used in aggregate for non-commercial academic research. We release code and aggregate statistics; we do not redistribute raw user comments beyond anonymized examples.

9. Reproducibility Statement

Code and data-processing scripts are publicly released. Settings: BGE-small-zh embeddings (512-dim, normalized); two-tower residual projections (512→128→512, dropout 0.3, weight decay $1e-3$, lr $5e-4$); early stopping on validation Recall@5 (patience 6); chapter-level train/val/test split (seed 42).

References

- [1] Zhang, W., et al. (2022). A Survey on Aspect-Based Sentiment Analysis. *IEEE TKDE*.
- [2] Zhao, J., et al. (2024). Sentiment Analysis of Video Danmakus. *Scientific Reports* 14.
- [3] Ouyang, L., et al. (2022). Training Language Models to Follow Instructions with Human Feedback. *NeurIPS*.
- [4] Radford, A., et al. (2021). Learning Transferable Visual Models from Natural Language Supervision (CLIP). *ICML*.
- [5] Xiao, S., et al. (2023). C-Pack: Packed Resources for General Chinese Embeddings (BGE). *arXiv:2309.07597*.
- [6] Reimers, N., & Gurevych, I. (2019). Sentence-BERT. *EMNLP*.
- [7] He, R., et al. (2017). An Unsupervised Neural Attention Model for Aspect Extraction.

ACL.

[8] Cohen, J. (1960). A Coefficient of Agreement for Nominal Scales. *Educ. Psych. Meas.*